

Supporting RISC-V Performance Counters Through Linux Performance Analysis Tools

João Mário Domingos

INESC-ID, Instituto Superior Técnico
University of Lisbon, Portugal
joao.mario@tecnico.ulisboa.pt

Tiago Rocha

INESC-ID, Instituto Superior Técnico
University of Lisbon, Portugal
tiagolopesrocha@tecnico.ulisboa.pt

Nuno Neves

INESC-ID, Instituto Superior Técnico
University of Lisbon, Portugal
nuno.neves@inesc-id.pt

Nuno Roma

INESC-ID, Instituto Superior Técnico
University of Lisbon, Portugal
nuno.roma@inesc-id.pt

Pedro Tomás

INESC-ID, Instituto Superior Técnico
University of Lisbon, Portugal
pedro.z.tomas@tecnico.ulisboa.pt

Leonel Sousa

INESC-ID, Instituto Superior Técnico
University of Lisbon, Portugal
las@inesc-id.pt

Outline

- Motivation
- RISC-V HPM
- Perf_events, Perf and PAPI
- Proposed approach
- Results
- Next steps

Performance Monitoring:

why it matters for system designers and software developers

There are multiple ways to improve software performance:

- Execution time profiling (count cycles)
- Hardware simulators (software-based, mixed)
- Hardware Performance Monitors (through Perf and PAPI)
- Mixed/proprietary tools (Intel Advisor, Arm Coresight)




Want to get faster?
Profile and tune your system

RISC-V HPM:

the RISC-V hardware performance monitor

Counters and Timers



v1.7 - 2015

Perf and PAPI

Perf_events Kernel Driver:

- Communication between user-privilege level and HPM
- Low-level; programmer must know encodings


Perf Application:

- Uses the perf_events
- Low-level
- Default events and raw events



PAPI API:

- Also uses perf_events
- Portability

RISC-V HPM + Linux Perf



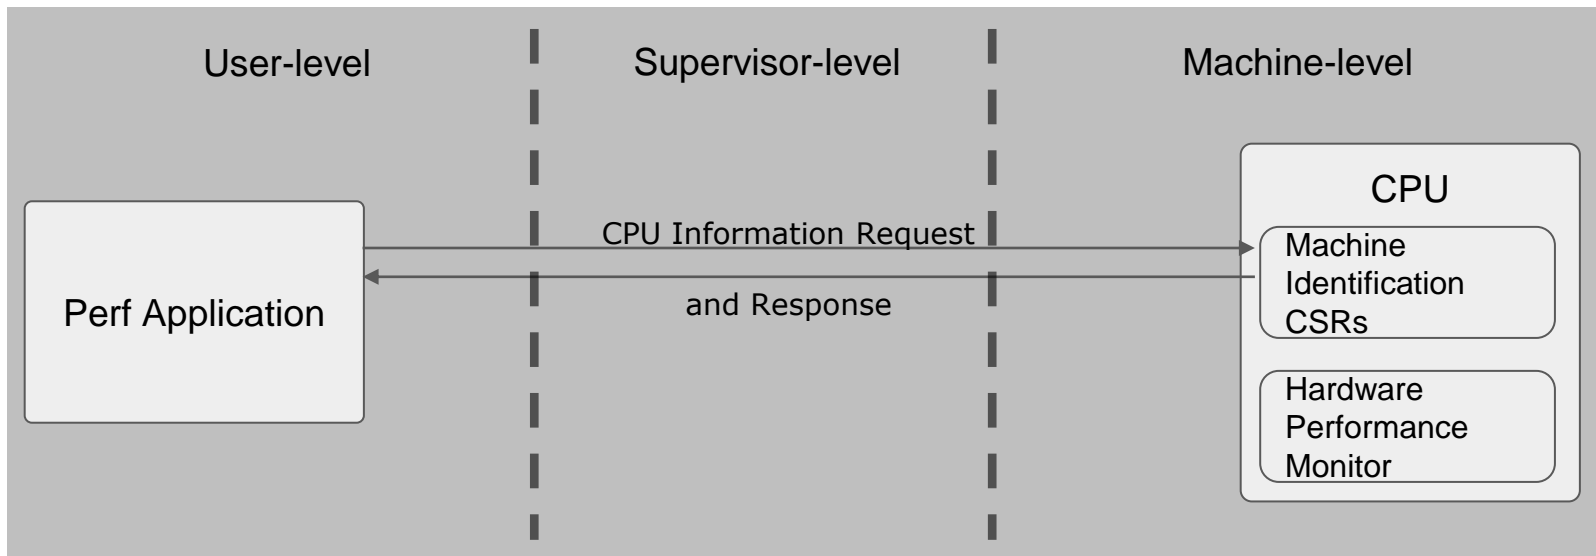
RISC-V HPM + Linux Perf: Software Overview



- CPU events identification
- Event attribution and counting (through the driver)
- Display results in an uniform way

RISC-V HPM + Linux Perf:

Software Overview



- **CPU events identification**


Contribution

- CPU PMU identification through the architecture ID and implementation ID
- Each CPU unique ID selects the appropriate set of events for that processor

RISC-V HPM + Linux Perf:

CPU specific event selection with Perf pmu_events

Contributions



CPU PMU Unique ID (composed by architecture and implementation IDs)

Used to match the CPU PMU with a set of event description code

CPU Identifier	File	Version	Events	Filename	Events	Type
0x300	, 0		CVA6		, core	
0x500	, 0		SPIKE		, core	
0x200	, 0		BOOM		, core	
...						

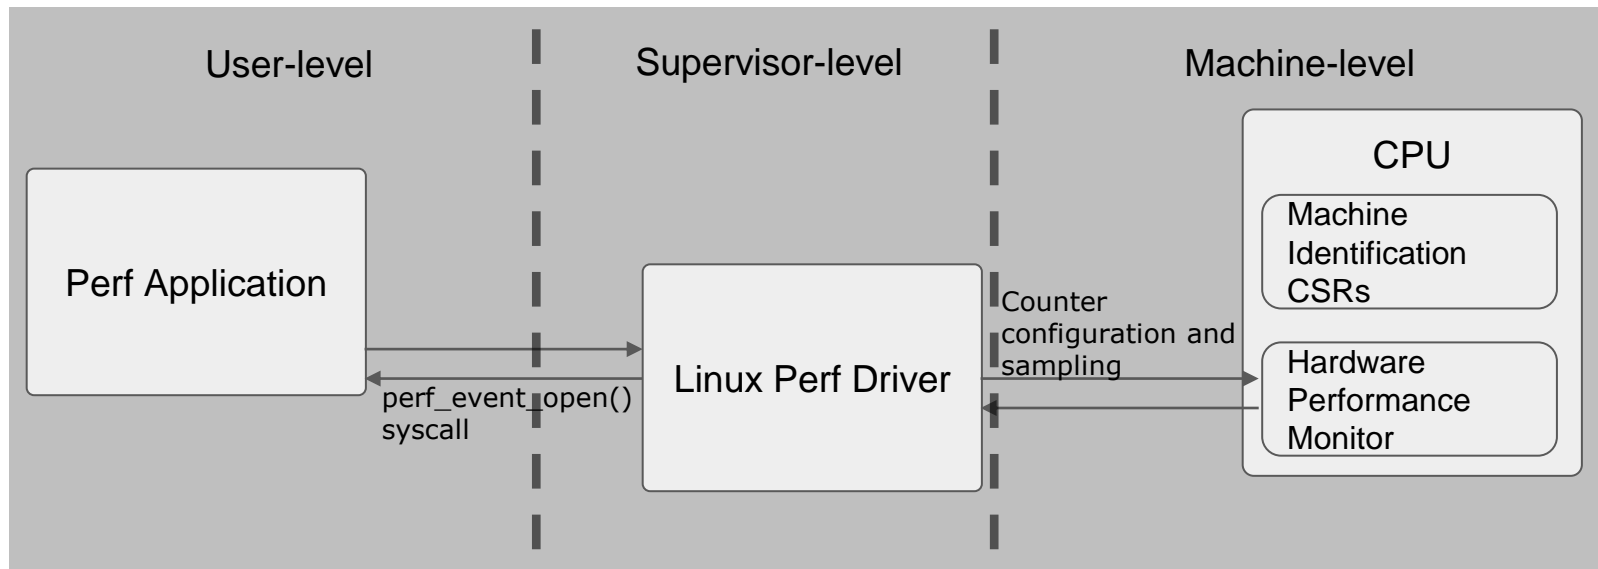
```
{  
  "Public Description": "This is an example event,  
  for demonstration purposes.",  
  "Brief Description": "This is an example event."  
  "Event Code" : "0x11",  
  "Counter Mask" : "0xF8FF",  
  "Event Name" : "EXAMPLE_EVENT",  
}
```

Events are simply described in a JSON format.

Each event is identified by the name, and provides an event code and a counter selection mask.

RISC-V HPM + Linux Perf:

Software Overview



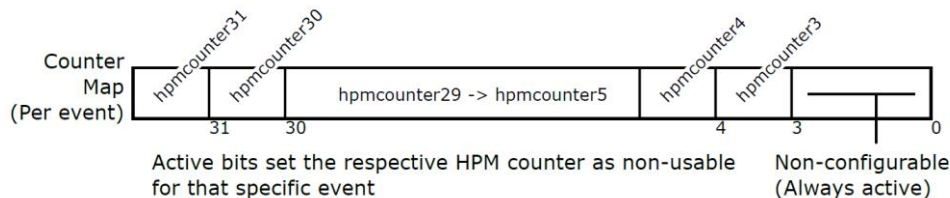
- **Event attribution and counting**

- Perf Initiates performance counting through a system call
- The Linux Perf driver setups events and samples counters

RISC-V HPM + Linux Perf:

Counting events with the perf kernel driver

Contributions

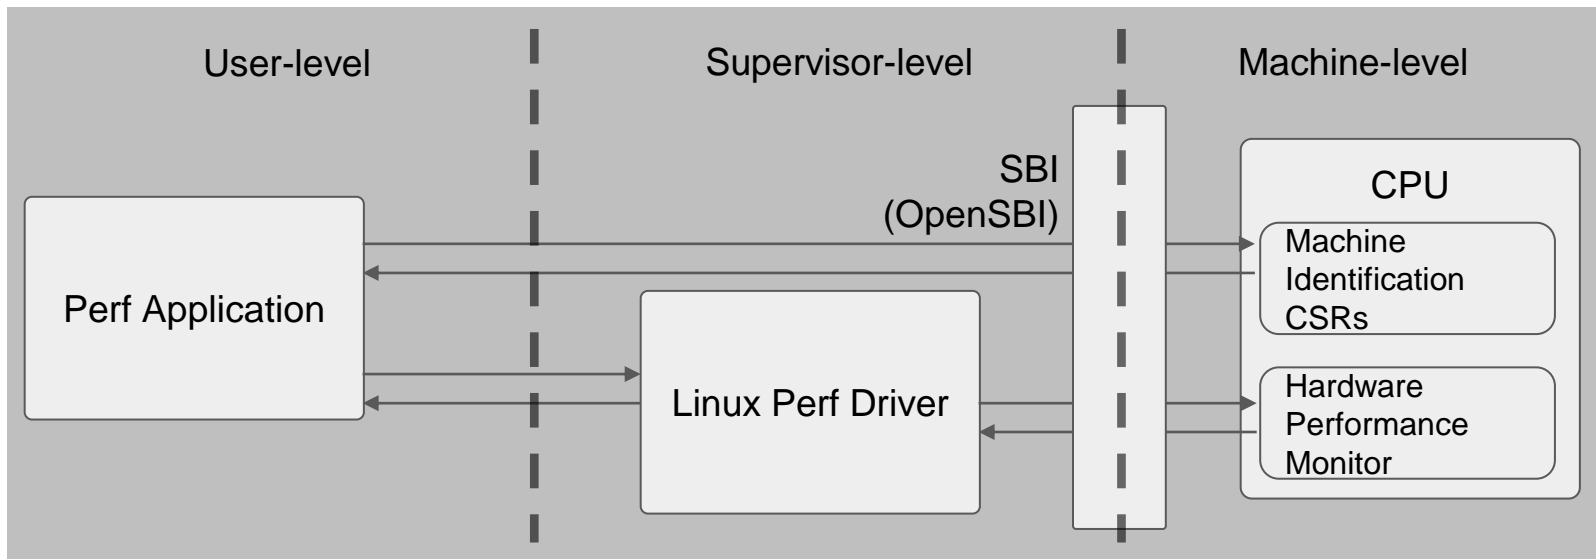


New perf event configuration allows for each event to have a counter mask, providing identification of available counters.

- Improved event configuration through the RISC-V *mhpmevent#* registers.
- Improved support for raw events

RISC-V HPM + Linux Perf:

Software Overview



- **Privilege escalation to machine-mode**

Contribution

- SBI/OpenSBI extension for HPM counters interaction (configuration, read, write)
- Added Perf and Perf Driver SBI calls for HPM extension and machine identification CSRs access

RISC-V HPM + Linux Perf:


Privileged access through OpenSBI

Contributions

RISC-V HPM Specification



SBI HPM Extension



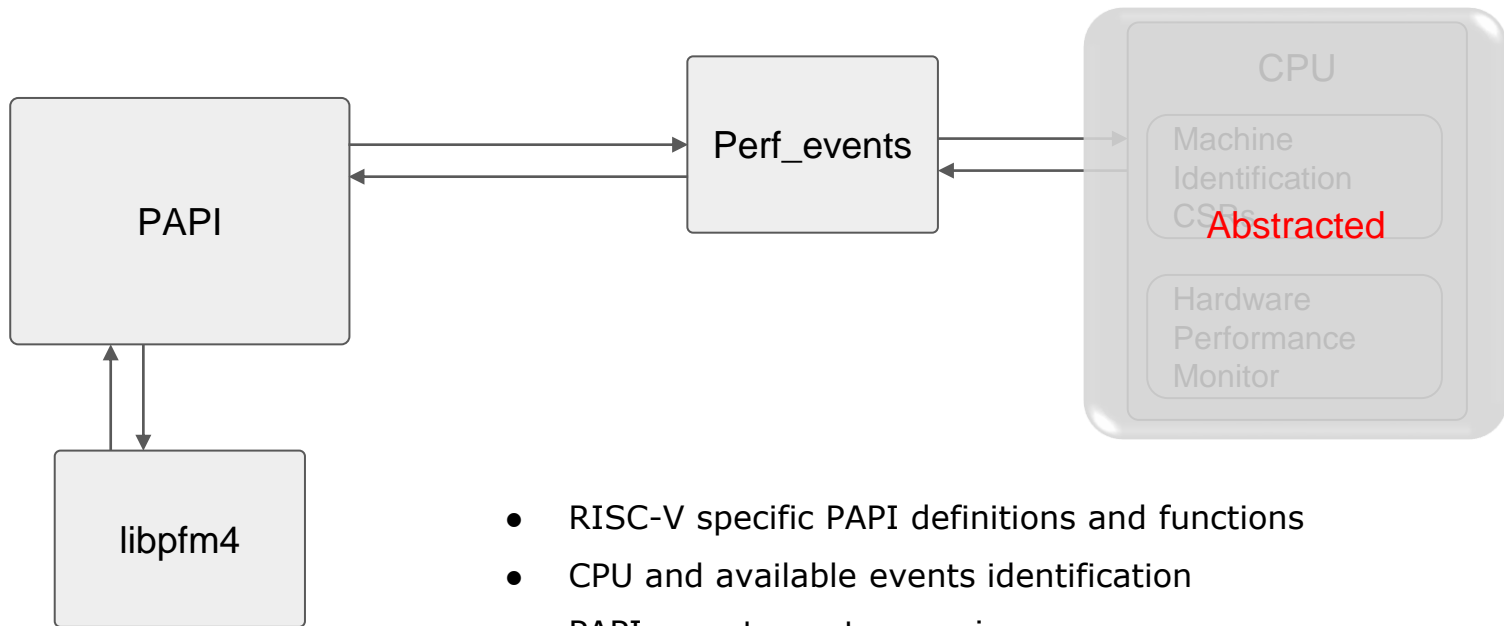
Current: v1.11 - 2019

RISC-V PAPI



PAPI (+libpfm4) + Perf_events:

Software Overview




- RISC-V specific PAPI definitions and functions
- CPU and available events identification
- PAPI preset events mapping

PAPI (+libpfm4) + Perf_events:

libpfm4 Changes

Contribution



- **Event decoding and event list manipulation**


- RISC-V event entry structure definition
- Functions to decode events and produce perf_event control structures and to iterate over event list
- Available events list
- HPM model description (#available counters, supported event list, ...)

```
{.name = "EXAMPLE_NAME",  
.code = 0x0000100,  
.desc = "Example description"}
```


PAPI (+libpfm4) + Perf_events:


PAPI Changes

Contribution



- **CPU identification and preset event mapping**
 - CPU identification through /proc/cpuinfo file
 - Inline assembly functions required by PAPI
 - PC context location definition
 - PAPI preset event mapping to available hardware events

Initial Results



Working Perf:

the *perf list* command

```
/ # perf list
branch-instructions OR branches      [Hardware event]
branch-misses                        [Hardware event]
cache-misses                         [Hardware event]
cache-references                     [Hardware event]
cpu-cycles OR cycles                 [Hardware event]
instructions                         [Hardware event]
alignment-faults                    [Software event]
bpf-output                           [Software event]
context-switches OR cs               [Software event]
cpu-clock                            [Software event]
cpu-migrations OR migrations         [Software event]
dummy                               [Software event]
emulation-faults                    [Software event]
major-faults                         [Software event]
minor-faults                        [Software event]
page-faults OR faults               [Software event]
task-clock                           [Software event]
duration_time                        [Tool event]
L1-dcache-load-misses                [Hardware cache event]
L1-dcache-loads                      [Hardware cache event]
L1-dcache-stores                     [Hardware cache event]
L1-icache-load-misses                [Hardware cache event]
branch-load-misses                   [Hardware cache event]
branch-loads                         [Hardware cache event]
dTLB-load-misses                     [Hardware cache event]
iTLB-load-misses                     [Hardware cache event]
```

CVA6 (Ariane) Events

Event	Counter	Event	Counter
Cycles	mcycle	Taken Exceptions	mhpmcounter9
Instructions Retired	minstret	Exceptions Returned	mhpmcounter10
ICache Misses	mhpmcounter3	Branches and Jumps	mhpmcounter11
DCache Misses	mhpmcounter4	Calls	mhpmcounter12
ITLB Misses	mhpmcounter5	Returns	mhpmcounter13
DTLB Misses	mhpmcounter6	Mispredicted Branches	mhpmcounter14
Loads	mhpmcounter7	Scoreboard Full	mhpmcounter15
Stores	mhpmcounter8	Instruction Fetch Empty	mhpmcounter16

Working Perf:

the *perf list* command

```
/ # perf list
branch-instructions OR branches [Hardware event]
branch-misses [Hardware event]
cache-misses [Hardware event]
cache-references [Hardware event]
cpu-cycles OR cycles [Hardware event]
instructions [Hardware event]
alignment-faults [Software event]
bpf-output [Software event]
context-switches OR cs [Software event]
cpu-clock [Software event]
cpu-migrations OR migrations [Software event]
dummy [Software event]
emulation-faults [Software event]
major-faults [Software event]
minor-faults [Software event]
page-faults OR faults [Software event]
task-clock [Software event]
duration_time [Tool event]
L1-dcache-load-misses [Hardware cache event]
L1-dcache-loads [Hardware cache event]
L1-dcache-stores [Hardware cache event]
L1-icache-load-misses [Hardware cache event]
branch-load-misses [Hardware cache event]
branch-loads [Hardware cache event]
dTLB-load-misses [Hardware cache event]
iTLB-load-misses [Hardware cache event]
```

Supports Perf default Hardware and Hardware Cache events

Not all Ariane Events are supported as Perf default events, we need **raw events**

CVA6 (Ariane) Events

Event	Counter	Event	Counter
Cycles	mcycle	Taken Exceptions	mhpcounter9
Instructions Retired	minstret	Exceptions Returned	mhpcounter10
ICache Misses	mhpcounter3	Branches and Jumps	mhpcounter11
DCache Misses	mhpcounter4	Calls	mhpcounter12
ITLB Misses	mhpcounter5	Returns	mhpcounter13
DTLB Misses	mhpcounter6	Mispredicted Branches	mhpcounter14
Loads	mhpcounter7	Scoreboard Full	mhpcounter15
Stores	mhpcounter8	Instruction Fetch Empty	mhpcounter16

Working Perf:

the *perf list* command

```
branch:
  ariane_branch_jump
    [Ariane branches/jumps count]
  ariane_call
    [Ariane calls count]
  ariane_mis_predict
    [Ariane mis-predicted branches count]
  ariane_ret
    [Ariane returns count]

cache:
  ariane_dtlb_miss
    [Ariane data TLB miss]
  ariane_itlb_miss
    [Ariane instruction TLB miss]
  ariane_l1_dcache_miss
    [Ariane data cache misses]
  ariane_l1_icache_miss
    [Ariane instruction cache misses]
  ariane_load
    [Ariane data loads]
  ariane_store
    [Ariane data loads]
```

Events Grouping Supported

Support for all CVA6 events

CVA6 (Ariane) Events

Event	Counter	Event	Counter
Cycles	mcycle	Taken Exceptions	mhpmcounter9
Instructions Retired	minstret	Exceptions Returned	mhpmcounter10
ICache Misses	mhpmcounter3	Branches and Jumps	mhpmcounter11
DCache Misses	mhpmcounter4	Calls	mhpmcounter12
ITLB Misses	mhpmcounter5	Returns	mhpmcounter13
DTLB Misses	mhpmcounter6	Mispredicted Branches	mhpmcounter14
Loads	mhpmcounter7	Scoreboard Full	mhpmcounter15
Stores	mhpmcounter8	Instruction Fetch Empty	mhpmcounter16

Working Perf:

the *perf* stat command

```
CoreMark 1.0 : 173.984251 / GCC10.2.0 -O2 -static / Heap
Performance counter stats for './coremark.linux.riscv':

2370362136      cpu-cycles
1467434059      instructions          #    0.62  insn per cycle
236016830      ariane_branch_jump
5313061         ariane_call
44041314        ariane_mis_predict
1406938         ariane_ret
1103           ariane_dtlb_miss
6870139         ariane_itlb_miss
2792754         ariane_ll_dcache_miss
8444217         ariane_ll_icache_miss
229115526       ariane_load
64636669       ariane_store
22573          ariane_exception
22573          ariane_exception_ret
239940589       ariane_if_empty
9386316         ariane_sb_full

23.804994040 seconds time elapsed

23.622965000 seconds user
0.109596000 seconds sys
```

CoreMark Result: 1.74 points/MHz @ 100MHz

Perf stat outputs event counts

There is support for event multiplexing,
but CVA6 does not have configurable events

CVA6 (Ariane) Events

Event	Counter	Event	Counter
Cycles	mcycle	Taken Exceptions	mhpmcounter9
Instructions Retired	minstret	Exceptions Returned	mhpmcounter10
ICache Misses	mhpmcounter3	Branches and Jumps	mhpmcounter11
DCache Misses	mhpmcounter4	Calls	mhpmcounter12
ITLB Misses	mhpmcounter5	Returns	mhpmcounter13
DTLB Misses	mhpmcounter6	Mispredicted Branches	mhpmcounter14
Loads	mhpmcounter7	Scoreboard Full	mhpmcounter15
Stores	mhpmcounter8	Instruction Fetch Empty	mhpmcounter16

Working PAPI (available components and preset events)

Available components and hardware information.

```
-----
PAPI version      : 7.0.1.0
Operating system  : Linux 5.13.19
Vendor string and code : RISC_V_SIFIVE (9, 0x9)
Model string and code : (0, 0x0)
CPU revision      : 0.000000
CPU Max MHz       : -1
CPU Min MHz       : -1
Total cores       : 4
SMT threads per core : 4
Cores per socket  : 1
Sockets           : 1
Cores per NUMA region : 4
NUMA regions      : 0
Running in a VM    : no
Number Hardware Counters : 4
Max Multiplex Counters : 384
Fast counter read (rdpmc): no
-----
```

Compiled-in components:

```
Name: perf_event      Linux perf_event CPU counters
Name: perf_event_uncore Linux perf_event CPU uncore and northbridge
      \-> Disabled: No uncore PMUs or events found
Name: sysdetect        System info detection component
```

Active components:

```
Name: perf_event      Linux perf_event CPU counters
                        Native: 204, Preset: 22, Counters: 4
                        PMUs supported: perf, perf_raw, riscv_sifive_u74
Name: sysdetect        System info detection component
                        Native: 0, Preset: 0, Counters: 0
```

PAPI Preset Events

Name	Code	Avail	Deriv	Description (Note)
PAPI_L1_DCM	0x80000000	Yes	No	Level 1 data cache misses
PAPI_L1_ICM	0x80000001	Yes	No	Level 1 instruction cache misses
PAPI_L2_DCM	0x80000002	No	No	Level 2 data cache misses
PAPI_L2_ICM	0x80000003	No	No	Level 2 instruction cache misses
PAPI_L3_DCM	0x80000004	No	No	Level 3 data cache misses
PAPI_L3_ICM	0x80000005	No	No	Level 3 instruction cache misses
PAPI_L1_TCM	0x80000006	Yes	Yes	Level 1 cache misses
PAPI_L2_TCM	0x80000007	No	No	Level 2 cache misses
PAPI_L3_TCM	0x80000008	No	No	Level 3 cache misses
PAPI_CA_SNP	0x80000009	No	No	Requests for a snoop
PAPI_CA_SHR	0x8000000a	No	No	Requests for exclusive access to shared cache line
PAPI_CA_CLN	0x8000000b	No	No	Requests for exclusive access to clean cache line
PAPI_CA_INV	0x8000000c	No	No	Requests for cache line invalidation
PAPI_CA_ITV	0x8000000d	No	No	Requests for cache line intervention
PAPI_L3_LDM	0x8000000e	No	No	Level 3 load misses
PAPI_L3_STM	0x8000000f	No	No	Level 3 store misses
PAPI_BRU_IDL	0x80000010	No	No	Cycles branch units are idle
PAPI_FXU_IDL	0x80000011	No	No	Cycles integer units are idle
PAPI_FPU_IDL	0x80000012	No	No	Cycles floating point units are idle
PAPI_LSU_IDL	0x80000013	No	No	Cycles load/store units are idle
PAPI_TLB_DM	0x80000014	Yes	No	Data translation lookaside buffer misses
PAPI_TLB_IM	0x80000015	Yes	No	Instruction translation lookaside buffer misses
PAPI_TLB_TL	0x80000016	Yes	Yes	Total translation lookaside buffer misses

Working PAPI (FP Stores test)

```
...  
float a, b, c;  
...  
for (int i = 0; i < 1000; i++) {  
    c = a*b;  
}
```

```
root@unmatched:~/tiagorocha/RISC-V-PAPI/sifive_u74_tests# ./papi_fp_store  
FP Stores: 1000
```


Next Steps...

- Ongoing: submitting perf_events kernel driver patch to the Linux Kernel repository
- Development of more tools that mimic the functionality of proprietary analyzes tools

This project was partially funded by

- FCT under projects UIDB/50021/2020
- European High Performance Computing Joint Undertaking under Framework Partnership Agreement No 800928 and Specific Grant Agreement No 101036168 (EPI SGA2)